

The Five-Level Assurance Ladder

Record, convention, advise, enforce, measure — how the toolkit keeps itself honest

1. What this answers

This document answers three questions you hit in the moment. A gate just refused something you did: what refused it, and why does it exist? You are adding a check of your own: which level does it belong at, and what does that level owe you? Someone asks whether the toolkit's assurances actually work: what is measured, and what is merely built?

The toolkit's assurances are not one mechanism. They are five levels, and each does a job the neighbouring levels cannot do. **Record** keeps facts alive past the session that produced them. **Convention** gives decisions named shapes. **Advise** names a miss without blocking anything. **Enforce** makes a silent omission impossible. **Measure** turns efficacy into a number instead of an assertion.

What follows takes them in order: section 2 is the law the whole ladder rests on, section 3 the five levels themselves, section 4 what each level buys that its neighbours cannot, section 5 the honest limit, section 6 one compact entry per shipped mechanism, section 7 the control surface, and section 8 how to add a level of your own.

A companion document, `MODEL_SUBSTITUTION_AND_VERIFIED_LAUNCH.machine.md` and its human twin, owns the serving-substitution failure and the patch built around it. This document cites that story where it bears; it never retells it.

2. The law beneath: structural over disciplinary

When you are choosing among candidate assurances, prefer a mechanism that fires at the decision moment over an exhortation that has to be remembered. An instruction written into a prompt fires only if the prompt was loaded and the model then chose to comply. A hook on the tool call fires regardless of both.

This is a conclusion drawn from recorded defects rather than a matter of taste, and three of the shipped mechanisms carry their own origin in their header comments.

The first is a hole in a model ban. The ban held at build time, where a barred identifier cannot ship in the payload, and at plan time, where a plan cannot declare one. A launch typed in the moment reached the barred model with no check at all, because neither of those gates sits on the dispatch path. The dispatch guard exists to close that hole and nothing else.

The second is a routing mandate that lived in the planner's prompt. It produced a plan declaring six delegation tracks, four of which the lead then ran itself. The plan-approval gate lints the routing block at the tool call, so the plan is corrected before the user ever sees it.

The third is a set of launch conventions for a verified frontier-tier child. They kept failing to reach sessions that had never loaded the documents stating them. The dispatch gate enforces the launch shape, so the convention arrives whether or not the doctrine did.

A corollary follows from the same reasoning. When an action's costs are asymmetric, default toward the cheap-to-reverse side and require an explicit token for the expensive one. Every gate described here fails open on its own internal error: it exits zero always, the refusal rides in structured output rather than in the exit code, and the error is logged. A broken guard must never wedge a session, and a silent fail-open is itself the bug, which is why the fail-open gets a log line.

The law says to prefer a mechanism. It does not say a mechanism can do everything, and section 5 states what no gate on this ladder can do.

3. The five levels

Level 1 — Record: facts survive

Nothing above this level means anything if the facts do not outlive the session that produced them.

Four **append-only ledgers** carry the record, each written so a worker with no prior context can read it. The change ledger takes one dated row per change and never edits a past row. The efficacy ledger takes one row per check, carrying its status and the measurement that justifies that status. The code inventory registers every script, so the next worker uses or adapts one rather than rebuilding it. The plan ledger records which plan is active, which is parked, and which is finished.

Receipts discipline governs how a current fact may be stated: only from a fresh read or measurement made for that claim. Recollection cannot be audited, and a command's output can.

Serving stamps decide the model identity of a run. That identity is read from the API response's own per-call `model` field in the run's own transcript, the only record downstream of what actually ran. Every other signal — a launch argument, a resolution, a header in the interface, the run's own testimony about itself — records intent or belief.

Completion telemetry appends one structured row per subagent completion, clean completions included, because a rate needs a denominator. Each row carries the hook build that wrote it, so a later analyst can separate pre-fix from post-fix populations instead of pooling them.

Level 2 — Convention: decisions have named shapes

A decision with no named shape cannot be checked by anything: not by a hook, not by a reviewer, not by the person who made it.

The project contract is a project-level `CLAUDE.md` stating the supervisory workflow, the folder contract, and the standing rules. It is installed into the project root rather than remembered.

The six-element brief fixes what a subagent gets: an assignment with a checkable done-condition; read-paths the child reads itself rather than receiving as a summary; a write-path inside the workspace for all products, code included; a report cap giving a line budget plus the receipts each claim must quote; the stop-when-stuck rule verbatim; and the scope rule verbatim. An empty element means the brief is not ready to launch. The fill-in form ships at `dev/briefs/_TEMPLATE.md`.

Routing blocks with owners require a plan's delegation block to name, for each track, the executor and also an owner that says more than the executor tag — the persona or skill assignment — or else to carry the explicit “no specialist fits” fallback.

The role assignment is the `ROLE:` line in a brief, naming which persona the child runs as, with the brief pointing at that persona's file.

Persona and skill delivery by read-pointer means the brief names the persona file or `SKILL.md` path and the child reads it. Two things follow. The child can audit what it was given, and for a frontier-tier child those opening reads double as the warmup that makes serving certifiable early.

Level 3 — Advise: misses get named, never blocked

This is the level for things that are usually an oversight and occasionally a considered choice, where refusing would be wrong more often than it would be right.

The **brief-slot advisory** names any unfilled brief slot at launch time and never refuses the launch. The **completion-telemetry nudge** puts one advisory note in front of the model when a finishing child's own final message scans at severity 2 or worse. The **watchdog scaffold** prints, after a frontier-tier launch returns, the ready-to-run certification command with the child's transcript path already substituted, together with

the exit-code legend, so the next action requires zero recall. **Plan-state re-injection** restores the active plan’s name, its snapshot path, and the resume protocol at each session boundary — startup, resume, and post-compaction — for roughly a hundred tokens per boundary event rather than a per-prompt tax.

Level 4 — Enforce: a silent omission becomes impossible

This is the level for omissions that are silent by nature, the ones nobody notices at the time, which is precisely why advice does not reach them.

The dispatch deny-gate refuses a frontier-tier launch that names no brief and opens with no warmup, or names a brief that is absent or unreadable, or names a brief missing its **ROLE:** line, its persona pointer, or, at the frontier tier, its warmup slot.

The plan-approval gate refuses the plan-approval tool call when the plan’s routing block is missing or malformed, owner-explicitness and multi-stage enumeration included, so the plan is fixed before the user is asked to approve it.

The turn-verdict block stops a turn once when that turn launched a certified-route child and names no serving certification, printing the ready-to-run certification command.

All of it is **one-shot by construction**: a turn-level block fires at most once per stop sequence, guarded by the harness’s own already-blocked flag. A gate that could fire twice is a loop, not a gate.

Level 5 — Measure: efficacy is a number, never an assertion

This is the level that decides whether any of the four below it actually worked.

Live certification reads a still-running child’s first serving stamps and returns a verdict by about the fifth call. **The post-hoc audit** catalogs every run’s launch argument against its resolved model against the model that was actually served, with a per-run verdict and a gate-able exit code. **Completion telemetry**, which Level 1 collects as rows, becomes a rate here: the rows carry denominators and build tags, so “did the fix work” is answerable by comparing populations rather than by assertion.

The adherence scorecard scores a session’s transcripts on five greppable dimensions as n over N , where n counts compliant cases and N counts the opportunities actually observed. It dates each opportunity by the transcript record that carries it, so an interrupted time-series across a gate’s ship instant stays interpretable instead of being smeared by sessions that straddle it. It measures; it never blocks and never edits. The workflow kit installs it at `<project>/dev/tools/adherence_scorecard.py`, with its fixture guard beside it, so the method by which the grades in section 6 were earned is one you re-run on your own sessions rather than merely a citation. Its lineage classification derives from your project root, so your sessions score as yours.

The efficacy-ledger discipline fixes the vocabulary every grade in section 6 uses, and it has five values. **attempted-untested** means built but not yet measured, and it is the status of anything new by default. **fixture-measured** means exercised against labeled fixtures or test cases, behaving as specified on those cases. **measured-working** means measured before and after on real outcomes, stated with the numbers. **measured-regressed** holds the same evidence standard and the opposite result. **retired** means no longer in use. No check moves past **attempted-untested** without a cited measurement, and “I built it” is never a measurement.

4. Why five — what each level buys that its neighbours cannot

Without the **record**, every other level is unauditable: a gate that fires and leaves no row cannot be told from a gate that never fired.

Without a **convention**, there is nothing to check. Levels 3 and 4 do not check goodness. They check conformance to a shape Level 2 defined, which is why every enforcement below points at a convention above.

Advice covers the large class where refusal would be wrong. An advisory that names a miss costs a line of context. A refusal that is wrong costs a retry, then a workaround, and eventually the operator switching the whole layer off, which is the real failure mode of over-enforcement.

Enforcement exists because advice reaches only an agent that reads it. The recorded failures in section 2 are all cases where the advice existed and did not arrive: the ban was written down, the routing mandate was in the prompt, the launch convention was documented. Enforcement is what fires when the doctrine did not load.

Without **measurement**, the other four are a story. **attempted-untested** is the default for a reason: a mechanism's existence is evidence about the author's intent, never about its effect.

That gives a rule for placing a new check. Choose the level by one question — what happens if the target behaviour is absent? If its absence is recoverable and often deliberate, advise. If its absence is silent and rarely deliberate, enforce. If you do not yet know, advise first, measure the rate, and promote afterwards. The tell to watch for: you are about to add an enforcement whose false-positive rate you cannot state, which means it should ship as advice and be measured.

5. The honest limit

Every gate on this ladder forces a decision to be **made**, **recorded**, and **well-formed**. None of them forces it to be **wise**. The owner-completeness and multi-stage checks are the clearest case: they are minimum-explicitness checks, and a routing block that names the wrong specialist for every track passes all of them.

The same boundary holds in general. These mechanisms check form and identity, not content-truth. A well-formed false claim passes. A brief with all six slots filled with nonsense passes. A certified child can be certified as the right model and still be wrong about the work.

The judgment residual is not left unowned. It belongs to **two-stage detection**. Every round that outputs a product a standard covers only partly runs two passes in order. First the automated detectors scan and disposition, with the results recorded. Then the standard's owner — the user by default, in a project with no style or standard document yet — sweeps the same output for what stage one missed. Output is never called clean on stage-one evidence alone.

Each stage-two miss then becomes three things: a labeled fixture pairing the flagged text with the corrected text, a row in the change ledger, and a routing decision, either to the next detector build or to the judgment checklist when no detector shape fits. Per-round residual counts accumulate, and the rate-drop trend across rounds is the real efficacy measurement. A single round's count is a baseline, not a verdict.

Every mechanism here can also be switched off, as section 7 describes. That is deliberate. An operator who sets the switch has made the layer inert on purpose, and a gate that ignored the switch would be a surprise. The consequence is worth stating plainly: these gates are a floor for a cooperating operator, never a control over a determined one.

Finally, a measurement is a statement about the conditions it was taken under. Where the measured thing is served remotely, its rates can drift on the scale of hours with no local change at all, so certification at launch carries the guarantee rather than configuration or a rate measured yesterday. `MODEL_SUBSTITUTION_AND_VERIFIED_LAUNCH.machine.md` holds that case in full.

6. Mechanism entries

Each entry below gives the file path an adopter has after install, where `<project>/` is the project root you installed the workflow kit into and `~/.claude/` is the global toolkit install; the harness event and matcher

it registers on; what triggers it; what you see when it fires; how to test it; how to silence it; and its evidence grade, copied from the efficacy ledger.

The contract every hook entry shares. Input is one JSON object on standard input. The exit code is always zero: a refusal rides in the structured output rather than in the exit code, because the model needs the reason in order to retry, and that retry is the point. Any internal error — no interpreter, unparseable input, an unexpected shape — is treated as indeterminate, so the hook fails open and logs. All of them are read-only apart from their own log lines.

The test every hook entry shares, and why it is valid. Every hook here is a pure reader of JSON on standard input, so replaying one crafted payload is the real firing path; nothing about a live session changes what the script reads. The form is `printf '<payload>' | bash <hook path>`, after which you read standard output — empty means a silent pass, one JSON object means a decision — and standard error, which carries advisory text. A worked example, a briefless frontier-tier launch, which must print a deny object:

```
printf '{"tool_name":"Agent","cwd":"/tmp/x","tool_input":{"subagent_type":"fable-executor"},"prompt":"do\n  | bash "<project>/dev/tools/test_enforcement_e2e.sh"'
```

Then invert it: point the prompt at a compliant brief and confirm silence. A gate you have only seen pass is not a gate you have tested. Show it refusing the thing it exists to refuse, then show it passing what it must not refuse.

Each hook also names its own per-hook regression battery in its header comment, as `tests/test_<name>.sh`. Those stay development-side and are not part of the release cut. The end-to-end battery, however, ships: `<project>/dev/tools/test_enforcement_e2e.sh` replays the whole violation set against the installed hook copies and against their `.claude/settings.json` registrations — the delivery layer, which a vendor tool-rename kills silently while every hook file stays perfectly correct. One command; exit 0 means every arm behaved. Arms that need a surface a given project may not have, such as the global toolkit install or a live plan under `plans/current_active/`, print SKIP rather than FAIL, because a false failure is how a check teaches its adopter to ignore it.

6.1 Advise

Brief-slot advisory — `<project>/dev/tools/test_enforcement_e2e.sh`, on PreToolUse with matcher `Task|Agent`. It triggers on a subagent launch: the prompt or description naming a `dev/briefs/*.md` file is looked up, a substantive launch naming none gets a one-line note, and a trivial launch (a read-only searcher, a short prompt, no write intent) passes in silence. What you see is injected context naming the unfilled slots, which are the six brief elements plus the `ROLE:` routing line, all advisory. It emits no permission decision at all: it advises, it does not deny. Silence it with `PLANNER_KIT_HOOKS=off`. Grade: `fixture-measured`.

Launch scaffold — `<project>/dev/tools/test_enforcement_e2e.sh`, on PostToolUse with matcher `Task|Agent`. It triggers when a certified-route or probe launch returns and the response carries an `output_file:` path; no path means it stays fully silent. What you see, on standard error, is the ready-to-run certification command with the child's transcript path substituted, plus the legend: FAITHFUL, exit 0, means certified, continue; SWAPPED at call k, exit 1, means relaunch or proceed knowingly and log it; UNDETERMINED, exit 2, means investigate. Standard output stays empty, since this hook decides nothing. Silence it with `PLANNER_KIT_HOOKS=off`. Grade: `fixture-measured`.

Plan-state injector — `<project>/dev/tools/test_enforcement_e2e.sh`, on SessionStart with matchers `startup`, `resume`, and `compact`. It triggers at a session boundary when a `plans/PLAN_LEDGER.machine.md` row under the project directory is marked ACTIVE; with multiple ACTIVE rows the last is named plus a count of the rest, and no ledger, no ACTIVE row, or any parse doubt leaves it silent. What you see is about two lines: the active plan's name and snapshot path, then the one-line resume protocol including the verified-launch reminder. Standard input is drained and ignored, and it writes nothing at all, not even a log, because a session-boundary hook is a pure reader. Silence it with `PLANNER_KIT_HOOKS=off`. Grade: `fixture-measured`, with one live output receipt against a real ledger.

6.2 Enforce

Dispatch deny-gate — `<project>/.claude/hooks/fable-dispatch-gate.sh`, on `PreToolUse` with matcher `Task|Agent`. This is the one deny-capable launch gate. It triggers on a launch whose subagent type is one of the two certified executor routes, or whose model argument names the ceiling tier; subagent types beginning `probe-` are allowlisted, since probe cells are legitimately briefless by design, and everything else passes silently.

It runs three checks. The first is launch shape: the prompt or description names a `dev/briefs/*.md` brief, or the prompt opens with the warmup token, and neither means deny, while a named brief that is absent or unreadable also means deny, because a dangling brief pointer is a launch defect rather than a parse error. The second is routing explicitness, applied once a brief resolves: the brief carries a `ROLE:` line and, anywhere in the file, either a persona-file reference or the literal fallback phrase, and a missing one means deny. The scope is the whole brief on purpose, because persona pointers legitimately ride the warmup slot rather than the `ROLE` line. The third is verified launch, at the ceiling tier only and once a brief resolves: the brief carries the warmup token, and a missing one means deny.

What you see is a denial of the launch with a reason of at most four lines: the check that failed, the element that is missing, the fix, and the template path. Silence it with `CRT_MODE=off` or `CRT_MODE=observe`; it is an intervention hook, so `on`, the default, is the only enforcing value. Grade: `live-measured`.

Plan-approval gate — `~/.claude/hooks/plan-routing-gate.sh`, on `PreToolUse` with matcher `ExitPlanMode`. It triggers when the agent is about to present a plan for approval, and it lints every routing fence and the markdown table, not just the first, so a compliant first block cannot shadow a violating second one.

The checks are the frozen routing-schema lint set. A declared delegation track must actually delegate, so body text saying the lead runs it is caught. Barred model aliases are denied wherever they appear, because an alias re-resolves silently and is therefore never a sanctioned route. A raw model field naming a restricted model is denied, while naming that model inside an executor, owner, task, or tradeoff field is allowed, so a plan that legitimately routes to a supervised executor can say so without the gate reading its own documentation as a violation. A supervised-executor track must carry its supervision marker, since the marker is what distinguishes a watched run from one that was routed and then abandoned. Owner-completeness requires a delegation track's owner to say more than the executor tag, or to carry the explicit fallback. Multi-stage completeness requires a track declaring a cascade topology to enumerate at least two stages, and a plan routing more than one delegation track to carry a waves or ordering section.

What you see is the approval call denied with the failing check named, so the plan is corrected before the user is asked to approve it.

One note on extraction. The current harness sends neither the plan text nor its path in the tool input, so the gate falls back to the newest plan file in the plans directory, accepted only when its modification time is within the last thirty minutes. Any miss — no directory, no file, too old, unreadable — is a logged fail-open rather than a lint failure. Silence it with `CRT_MODE=off` or `observe`. Grade: `live-measured`.

Alias dispatch guard — `~/.claude/hooks/opus-dispatch-guard.sh`, on `PreToolUse` with matcher `Task|Agent`. It triggers on a launch whose model argument is a bare restricted alias, alone or context-suffixed; everything else passes silently, whether that is a named tier alias, a full model identifier, or no model key at all.

Why aliases and only aliases: an alias reads to a human like a tier choice while resolving at launch to whatever the launcher currently maps it to, and it re-resolves silently whenever that mapping changes. A full identifier is deliberately not denied here, because the launch parameter accepts aliases only, so a full identifier is rejected upstream by the harness, and because the sanctioned supervised route names an agent whose own frontmatter carries the pin rather than naming a model.

What you see is the launch denied with the structured reason. Silence it with `CRT_MODE=off` or `observe`. Grade: `fixture-measured`, from matrices with no-guard red baselines; the code inventory additionally records live firing in the session it shipped.

Turn-verdict block — `<project>/.claude/hooks/collect_outcome_gate.sh`, on `Stop`. It runs two layered checks over the segment after the last genuine user message. The outcome check fires when the turn collected subagent results and the reply names none of the six collect outcomes; the signal that the turn saw subagent results is read from the transcript by three independent tells, any one of which suffices, and no tell means silence, so the gate never fires on a turn it cannot read. The verdict check fires when the same segment launched a certified-route child and carries no certification token, meaning either a verdict word or a serving-stamp receipt line. That token is searched in decoded content text only, so a transcript record's own model field can never masquerade as a receipt, and a user prompt, which the segment excludes, can never certify a child it merely mentions.

What you see is one block carrying the request. A `Stop` hook has no advisory channel: blocking is the only way to put text in front of the model, and it costs one extra turn, which is why the trigger is narrow. On the verdict check the block names the ready-to-run certification command with the child transcript path substituted from the launch result. The harness's already-blocked flag covers both checks, so there is at most one block per stop sequence and never two. Silence it with `PLANNER_KIT_HOOKS=off`. Grade: `fixture-measured` for the outcome check and `fixture-measured` for the verdict check.

6.3 Measure

Live certification — `<project>/.claude/skills/model-verification/fable_watchdog.py`, also installed globally with the toolkit's skills. Run it as `python3 <path>/fable_watchdog.py <child-transcript> [--expect <model-id>] [--verdict-at 5] [--watch]`. You trigger it, at or just after a launch, and `--watch` polls a still-growing transcript, so it can be pointed at a running child. What you see is one verdict line, and the exit code is the verdict: 0 for `FAITHFUL`, meaning the first `--verdict-at` stamps all equal `--expect`; 1 for `SWAPPED` at call `k`, where the first divergent stamp is call `k`, decided the moment a divergence is seen; and 2 for `UNDETERMINED`, meaning too few stamps yet, no assistant records, or unparseable input, and never a crash. It pairs with the warmup convention: the brief's persona and skill read-pointers are the opening calls, so the transcript reaches a certifiable length through useful work rather than throwaway calls. Grade: `live-measured`.

Post-hoc audit — `~/.claude/skills/model-verification/model_run_audit.py`. Run it as `python3 <path>/model_run_audit.py <project-dir | session.jsonl | session-dir> [--json|--tsv|--summary-only]`. What you see is every run catalogued across its record layers — the raw launch argument, the harness's resolution, and the API's own served stamp — with a per-run verdict of `MATCH`, `MISMATCH`, or `UNKNOWN`, plus totals by served model and the mismatch signature. The exit code is 0 for no mismatch, 1 for at least one mismatch, and 2 for an input or parse error; `UNKNOWN` never gates, because a main-loop run carries no intent layer by construction. One caution: a naive `grep` for `"model"` over a transcript that launches children also matches each launch's model argument and config echoes, mixing intent into what reads as a serving tally. The audit script separates the layers, and the skill documents an accurate one-liner for when the script is unavailable. Grade: `fixture-measured`, plus one live acceptance run.

Completion telemetry — `~/.claude/hooks/subagent_qa_gate.sh`, on `SubagentStop`. It triggers when a subagent finishes, writing one structured row per completion, clean completions included, and putting an advisory nudge in front of the model only when the scan found something at severity 2 or worse.

The scanned text is the finishing child's own final message, taken from the payload the harness supplies where present and otherwise resolved through a documented ladder of narrowing fallbacks, with each row recording which resolution was used. This matters because the transcript path supplied on this event is the parent session's file, so "the last assistant message in the transcript" is usually the coordinator's text rather than the child's, a measured source of false flags, each of which cost a turn to refute.

One of its findings is not text-shaped but value-gated: a certified route's serving stamp is compared against that route's shipped model promise, so a silent substitution becomes loud with zero operator discipline required. What you see is a row appended to the error-mode log, and a nudge only above the severity threshold. Grade: the child-scoped scan is `fixture-measured` with a live signal on real rows, and the substitution flag is `fixture-measured`, since no real substitution has crossed it live yet.

7. The control surface

There are two switches, split by what a mechanism does rather than by which file it lives in.

`PLANNER_KIT_HOOKS=off` silences the advisory set: the brief advisory, the launch scaffold, the plan-state injector, and the turn-level Stop hook. These inject context or block a turn, so they must be silenceable.

`CRT_MODE` gates the intervention set: the dispatch deny-gate, the plan-approval gate, and the alias guard. The value `on`, which is the default, enforces; `off` and `observe` both silence. It is set from the environment, or from the mode file the toolkit reads when the variable is unset.

Silencing is an operator act with a whole-layer effect, not a per-case escape. When a gate refuses something legitimate, fix the shape it is checking — fill the slot, name the owner, add the warmup — rather than switching the layer off. The refusal names the missing element and the template precisely so that the fix is cheaper than the bypass.

8. Adding to the ladder

Six steps. Name the failure the check exists to catch, and the record that shows it happened; a check with no recorded failure behind it is a guess. Pick the level by the one question in section 4. Write the check so that it fires at the decision moment, fails open on its own error, and logs that fail-open. Show it red before green: plant the exact fault it exists to stop and confirm it fires, because an untested gate is an unverified claim one layer down. Register it at `attempted-untested` and say so where it ships. Promote it only on a cited measurement, and state the population that measurement covers.

The shape generalizes past this toolkit. Any system where an instruction can go unread by the party it binds meets the same split: a durable record, a named shape, a cheap advisory for recoverable misses, a hard gate for silent ones, and a measurement that decides whether the previous four earned their keep. What does not transfer is any particular rate, because a rate belongs to the conditions it was measured under.

The record outlives the session; the convention gives the check something to check; advice covers what refusal would get wrong; enforcement covers what advice cannot reach; and measurement decides whether any of it worked. No gate anywhere on the ladder claims to make a decision wise, which is why the human sweep at stage two is part of the architecture rather than an admission about it.