

Inside the Loop

How the supervised workflow runs, and how you steer it

You asked for help improving your scientific writing. Behind that one request, a small team assembled, did its work, and disbanded, several times over. A coordinator broke the job into pieces, handed each piece to a specialist, waited for the results, decided what to do next, and only then assembled the next team. This guide explains that loop: what it is, why it is built the way it is, and how to predict what it will do next and steer it when you want to.

The work runs this way for a concrete reason. A single assistant, working alone, has one **context window**: one finite pool of working memory that holds everything it is currently thinking about. Your writing-improvement job does not fit in that pool. It spans a corpus of sixty curated failure examples pulled from a PDF, a body of detector code that flags weak prose, a benchmark that measures how well those detectors work, and a rendering pipeline that turns machine notes into a readable document.

Ask one assistant to hold all of that at once and one of two bad things happens. Either it loses the thread, re-reading the same files until it starts to contradict itself, or it fakes competence, claiming a fix works without ever measuring whether it does. The supervised loop exists to prevent both failures. By the end of this guide you should be able to look at any moment in the work, say what the loop is doing and why, and step in where your input changes what happens next.

What it is

The workflow is a **supervised loop** with four named parts. Three of them do the work.

A single coordinator, the **planner**, decomposes the goal, routes each piece of work, launches the workers, collects what comes back, and decides the next move; apart from a few narrow exceptions, it does no domain work itself. A rotating cast of **subagents** does the reading, writing, building, and measuring. And **custom code and software** handles the invariant, deterministic tasks: the work the project has to repeat, and the work that accuracy or completeness makes safer in code than in prose. Code sits at the same level of importance as the planner and the subagents, because it is one of the three things that actually does the job.

The fourth part is the **durable files** in a shared folder tree. They are what the other three produce and pass between one another: briefs, reports, ledgers, measurements, and the code itself.

A subagent is not the same thing as the assistant you are talking to. A subagent is a **fresh, separate instance** of the assistant, launched with a single written assignment and no memory of the wider conversation. It cannot see what you said, what the planner is thinking, or what its sibling subagents are doing. It sees exactly what its assignment tells it to read, it does its one job, it writes its output to a named file, and it returns a short report. Then it is gone. This isolation is deliberate.

The picture below is the loop as one cycle: the planner at the top, the two kinds of worker below it, and the shared file tree at the bottom. One arrow runs back up: what the workers write is what the planner reads to decide the next cycle.




Figure 1: The supervised loop: the planner, its two kinds of worker, and the durable files that carry everything between them.

The planner sends each subagent a **brief**: an assignment, the exact files to read, and the file to write. It also chooses their arrangement: one worker, several at once, or a chain. Deterministic work goes to code instead of to a subagent’s narration, and that code runs inside a subagent by default, or directly under the planner when a quick check is all it needs. Whatever any of them produces lands in the shared tree as a durable output. The planner reads those outputs, decides whether the plan still fits, and launches the next cycle.

One caveat about the picture. Its arrows can read as live steering, as if the planner were leaning over a subagent’s shoulder while it works. Mostly it is not. Each arrow is a **per-cycle** exchange: a launched subagent does its job and returns one report, and the planner’s ordinary way of correcting course is to wait, collect, and launch a fresh subagent with a corrected brief. One narrower channel does exist. The planner can send a message to a running subagent, and the message reaches it at its next tool call, so a correction or a stop request can land while the work is under way. What that channel cannot do is redesign the job: a subagent works from its brief, and a new design means a new brief.

An example from your project

To make this concrete, here is a stretch of the supervised loop that ran in your project. It ran as two cycles, one after the other.

The goal of one early phase was to sharpen the writing detectors using your own curated examples of bad scientific prose as the yardstick. The planner broke that goal into pieces and ran the first piece as a single cycle. It briefed one subagent to **extract** the examples from your PDF, `WRITING_FAILURE_EXAMPLES_NP.pdf`, into a structured machine file. The subagent read the PDF, pulled sixty verbatim example atoms, sorted them into thirty of your failure classes, built a table mapping each class to the detector that ought to catch it, and verified its own counts before reporting. Its durable output was one file, `dev/failure_examples_extract.machine.md`. The planner collected that output, read the report, and recorded the change in a running log. That is one **iterative cycle**: brief, launch, durable output, collect, log.

The next cycle built on it. Now that the sixty examples existed as ground truth, the planner routed a second, different subagent, a software specialist this time, to build the **yardstick itself**: a small harness, living in `dev/benchmark/`, that runs every one of the sixty examples through the current detectors and reports which are caught and which slip through. The harness is code; running it produces a measurement; the measurement tells the project where its detectors are weak. The first cycle produced a durable file, and the second cycle read that file and produced a measurement. The output of one cycle became the input of the next, and the two subagents never communicated at all; their only connection was a named file in the shared tree.

You are holding another instance of this pattern. This guide was not written by the planner. The planner routed the job as a two-stage chain: a first subagent, a rationale analyst, read your diagram and notes and distilled the workflow's reasoning into a machine-readable ledger, `dev/workflow_rationale_ledger.machine.md`; then a second subagent, the writing specialist that authored these words, read that ledger and wrote the guide you are reading. The writing specialist never saw the analyst's conversation. It saw one file.

The iterative cycle in more detail

The **iterative cycle** is the workflow's unit of work: everything the loop does happens inside one cycle or another. A cycle is one full pass, from plan through brief, launch, collect, and synthesis, and it ends in a decision. Each cycle closes by asking a single question: given what just came back, does the plan still fit the evidence? That question is what keeps the work honest. The plan is never executed blindly to the end; it is re-fitted to fresh evidence at the close of every cycle.




Figure 2: The iterative cycle: one pass, from plan to the collect-time decision.

The **brief** is where a cycle succeeds or quietly fails, so it carries a fixed checklist of six elements, and an element left empty means the brief is not ready to launch. First comes the **assignment**, with a done-condition a reader can check against an artifact rather than against the subagent's word. Then the **read-paths**, the sources the subagent reads for itself, because a summary pre-decides what mattered and a subagent cannot audit what it never saw. A **write-path** follows, one in-workspace destination for every work product, code included: a subagent given none falls back to its own temporary scratch space, and its work evaporates when it ends.

The last three elements govern how the subagent comes back. A **report cap** sets a line budget and names the receipts each claim must quote, since the durable file is the deliverable and the report only points at it. The **stuck rule** and the **scope rule** are both handed over word for word, the one so a bounded subagent halts and reports instead of thrashing, the other because a fresh subagent inherits nothing, and a boundary held only in the planner's context never reaches the party that would breach it.

Naming the six does something a prose description cannot: written as duties, the anatomy can be satisfied in spirit while one element quietly goes missing, but numbered, with a fill-in form at `dev/briefs/_TEMPLATE.md` that the planner copies and fills in full, an unfilled element becomes an object someone can point at. Each non-trivial brief is then persisted to `dev/briefs/` before launch, so a crashed session re-launches from the file instead of from a lost conversation.

The diamond has six exits, and they are the planner's entire repertoire of moves. Reaching it takes two acts, not one: the planner first spot-checks the receipts in each report against the artifacts they name, because a claim in a report is not the artifact, and only then asks the fit question outright.

The answer goes into the change log as exactly one of six named outcomes, along with the receipts it checked. **CONTINUE** means the plan holds and the next wave goes out. **RE-ROUTE** re-briefs only the pieces the evidence touched. **FIX-FIRST** says a defect blocks the next wave and has to be repaired before anything else proceeds. **ABORT** stops the work and reports the blockage to you. **GOAL-MET** is terminal, and it binds two acts: update the ledgers and take the dated backup. **ADAPT** means the evidence moved the target, so the planner breaks the goal into different pieces. Every cycle lands on exactly one of these, and the naming is the point: an unnamed outcome defaults silently to **CONTINUE**, which is how a plan outlives the evidence that justified it.

The arrangement of subagents within a cycle is a deliberate choice the planner makes each time, based on how the pieces of work depend on each other. A **single-thread** cycle is one subagent doing one job. A **sequential build** chains subagents so that each depends on the last; the extraction-then-benchmark pair above, and the analyst-then-writer chain that produced this guide, are both sequential builds. A **verify-loop** pairs a builder with a separate, fresh-eyed reviewer, because a builder is a poor reviewer of its own work.

The remaining arrangement is the fan-out. A **wave** is a launch in which the planner sends several subagents out at once, and it comes in two forms that look identical at launch. In a **parallel wave**, each subagent gets a different piece of the work and its result stands on its own; the planner collects the results and moves on. In a **convergence**, the planner gives several subagents the same material but a different question each, then merges their answers into one. The launch looks the same in both. What separates them is what happens at collect: a parallel wave needs only collection, while a convergence needs the output to be merged and synthesized, which is what the planner does.

When a piece of work earns its own planner

A piece of work earns its own planner when three things hold at once: it is internally multi-step and delegable, so it would need a plan of its own; it is separable behind a narrow interface of named inputs and outputs; and it is heavy enough that watching its internals would consume the planner's attention at collect. Then the piece goes whole to a single subagent running the planner role at the top tier, and that **sub-planner** runs the entire cycle described above inside it. The reason is plain: a piece that needs its own plan either gets one from a subagent that owns it, or it fragments across the planner's collect.

A sub-planner is sealed. Its brief carries the same six elements, with the scope narrowed to the subtree it owns, its read-paths set to the named inputs, and its write-path set to that subtree plus a private brief area, `dev/briefs/<ID>/`, for the subagents it launches. That private area exists because the shared ledgers have exactly one writer, the planner, and two writers appending to one running log lose or interleave each other's rows. Where several sub-planners run at once, the planner gives each a write scope that overlaps no other's, since a shared tree has no lock. What comes back is a **roll-up** rather than a transcript: within its report cap the sub-planner reports the outcomes it named and the receipts it checked, not the reports of the subagents under it, and the planner spot-checks those receipts and writes the shared row itself.

One limit was measured rather than assumed. Before the option was written down, a subagent running the planner role launched a subagent of its own and got the reply back, which is what lets a sub-planner supervise its own waves at all. The test covered one level, so the subagents a sub-planner launches are workers, not a third tier of planners.

The picture below puts the two routes side by side: the ordinary one along the left, and the sealed sub-scope where a sub-planner runs the same cycle over its own workers and returns a roll-up.




Figure 3: The sub-planner option: a chunk that earns its own plan goes whole to a subagent running the planner role, which runs the same cycle inside a sealed sub-scope and returns a compressed roll-up.

Why it works

Isolation is why durable files are the only channel. A subagent cannot see the planner's conversation or its siblings' work. So anything that must travel *between* agents has to travel through something they can all reach, and the only such thing is a named file in the shared tree. An unwritten result is invisible: it dies with the subagent that held it. This is why the workflow is fastidious about write-paths, and why the shared tree is not merely storage but the communication channel itself. When the analyst wrote the rationale ledger and vanished, that file was everything the writing specialist received.

Token economy is why the models are tiered. Every unit of an assistant's working memory and every unit of the work it does costs **tokens**, the currency of both memory and labor. A more capable model costs more per token. Spending a top-tier model on mechanical drudgery wastes

money, and spending a cheap one on subtle reasoning risks getting it wrong, so the planner **routes by difficulty**. The hardest reasoning goes to a named specialist running the top model; other hard reading, design, and code go to the tier below; medium synthesis, review, and sweeps go to a mid-tier model; purely mechanical fan-out goes to the cheapest. Routing asks its questions in a fixed order, and the picture below puts them in that order: first whether the task should run as code at all, then who runs it, and only then which model the rest of the work needs.




Figure 4: Routing a task: run it as code or delegate it, then decide who runs the code and which model the work needs.

The tier ceiling in that diagram is a firm rule: subagents run at the project’s top model *or lower*, and **never** above it, never the next-generation model, and never an *unpinned* model name. The top model is `claude-fable-5`, always available for the subagent tasks that most need it: higher-order reasoning, careful logic, high complexity, or many issues tracked at once. That availability holds regardless of the planner’s own level. The planner leads by position rather than by capability: it briefs, collects, decides, and checks the outputs that come back, so a planner at any level can still send the hardest piece to a top-tier subagent.

You barred both names for a specific reason: you found that Opus 5 makes reasoning errors and jumps ahead of the evidence much more than Opus 4.8, so you locked Opus use to Opus 4.8 at maximum reasoning. You banned the unpinned name for the same reason: a name with no version attached can silently resolve to whatever the newest model happens to be, which could be the model you excluded, and pinning the exact version removes that possibility. The floor and ceiling together keep every subagent inside a known, allowed band of capability and cost.

Determinism is why code is a primary output. The first question the planner asks of any task is whether it can be expressed as a fixed sequence of commands against code that already exists. If it can, the planner does not hand it to a subagent to narrate its way through; it has the work **run as code**. This is the *instrument test*, and it is the gate at the top of the routing diagram. Running the benchmark harness over the sixty examples is deterministic: the same inputs give the same outputs, every time, cheaply and exactly. Ask a language model to “act like” that harness and you get something slower, costlier, and non-reproducible, an approximation where you wanted invariant instrumentation or measurement.

Where that code runs is a second, separate question. A subagent runs it by default, which keeps the planner to coordinating; the planner runs code itself only for a quick verification or a one-command check. Either way the rule is plain: deterministic work is run, judgment work is delegated, and code is both one of the three things that does the work and, often, a primary output of it, as central as any written analysis.

So a subagent never discards the code it writes. Every script lands in the shared tree, and a **code inventory** records what each one does. Every brief then tells the subagent to read that inventory before writing anything new: use an existing script where it fits, extend one where it is close, and build from scratch only where neither works. A subagent that cannot see what already exists will build it again. That wastes time, wastes tokens, and sometimes introduces errors that take the work off its goal.

A summary can be lossy, so briefs carry read-paths. When the planner briefs a subagent, it does not paraphrase the source material and hand over the paraphrase. It names the **exact files** the subagent must read for itself. The reason is a quiet but serious hazard: a planner’s summary is a proxy for the real thing, and a proxy can be lossy or subtly wrong. A subagent that acts on the planner’s summary inherits the planner’s errors and blind spots, and compounds them. A subagent that reads the primary source can catch what the summary missed. This is why the benchmark builder was told to read the extraction file directly, and why the writing specialist was pointed at the rationale ledger itself. Reading the source is how a subagent minimizes the risk of inheriting any mistakes that the planner makes.

Efficacy-from-existence is why nothing “works” until it is measured. This kind of work has a common failure: building a fix and, because it now *exists*, believing it *works*. Writing a detector does not show that it catches anything. The workflow guards against this with an **efficacy ledger** whose governing rule is blunt: no check’s status moves past “attempted but untested” without a cited measurement. This is why the benchmark cycle existed at all. The sixty-example yardstick is the *measurement* that lets a detector’s status graduate from “written” to “verified.” Alongside it runs a second ledger, an append-only log of every change made in the workspace, the provenance trail that records the reason for each one. You add new entries to that log and never rewrite old ones.

A subagent’s near-independence is why most corrections wait for collection time. A launched subagent runs to completion and returns a single report. A message sent to it mid-run does arrive, at its next tool call, so a correction or a stop request can reach work already under

way. A redesign cannot: the brief is what the subagent is working from, and a different design is a different brief. So when the planner learns something mid-cycle, small corrections go out immediately and anything larger waits for the collect, where a fresh subagent is briefed for the pieces of work the new lesson touches. Background shell commands are different again, and can be adjusted freely while they run. For you, this means a correction lands quickly and a change of direction lands at the next collect.

A wrong approach is cheapest to catch early, so every brief carries a stuck rule. Every brief ends with the same instruction, handed to every subagent verbatim: *if errors recur, the approach stops converging, or you are about to change approach, stop and report back with what you found; do not thrash.* The rule does two things. It keeps a subagent from burning tokens grinding toward a wrong answer, and it turns a confused worker into a source of information, because what it found before it stopped is evidence the planner can act on.

The rule has already paid off here. In one cycle of your project, a subagent was dispatched to inventory a set of document files, discovered that the premise of its task conflicted with what it found, and **stopped and declined to build**, reporting the conflict instead of forcing a result. No wrong file was written. The planner read the report, dropped that line of work, and moved on. That is the stuck rule and the collect-time decision working as designed: stopping when the task is wrong is more useful than producing something wrong.

The guardrails the picture leaves out

The loop diagram at the top of this guide shows the *shape* of the work: who talks to whom, in what cycle. What it cannot show is the set of **guardrails** that keep the work safe and honest, and those guardrails are as much a part of the workflow as the loop itself.

The first guardrail is the **scope rule**, and it is absolute. All work, the planner's and every subagent's alike, stays *inside* this project's workspace folder and its subfolders. No subagent reads or writes anything outside it without an explicit, per-excursion grant from you. This rule is repeated, word for word, in every single brief, so that a fresh subagent with no memory of the conversation still inherits it. It exists because of a real incident: earlier reconnaissance subagents, left unbounded, went wandering into parent and system folders they had no business in. The rule was issued to ensure that never recurs. Isolating this sandbox is the entire point of having one.

The second guardrail is the **model floor and ceiling** already described: a hard band on which models a subagent may run, never above the project's pinned top model and never an unpinned name. The third is the **planner's reading discipline**, which has exactly two exceptions.

By default the planner does **not** read project files or code directly; that reading is delegated, which serves both token economy and the clean separation between a coordinator and its workers. The planner reads directly in only two situations. One is routine: when the thing to read is a subagent's output or report, which is its whole job to synthesize. The other is the accuracy backstop, invoked when a decision is too important to rest on an unverified claim, either to verify a load-bearing claim before a consequential decision or to resolve contradictory evidence at its source. Everything else, the planner delegates. Running code follows the same spirit: a subagent

runs it by default, and the planner runs it itself only for quick verifications and one-command instrument steps.

Two more guardrails came up in the worked example. The **two ledgers** are the efficacy ledger, which forbids the word “works” without a measurement, and the append-only change log, which records the reasoning behind each change. The **backup convention** covers the rest: a dated zip archive taken at a safe point, so you can undo a change wholesale. It sits alongside version control, not in place of it. Every brief also carries the six elements taught above, three of which are guardrails in their own right: the **read-paths** that point a subagent at primary sources, the **report cap** that bounds what it hands back, and the **stuck rule** that tells it when to stop. None of these appear in a loop diagram, but they are part of the workflow all the same.

These guardrails are carried at four strengths, each rung firing at a different moment. Weakest is memory, the planner’s own durable notes, which it may or may not consult. Above it sits prose, the standing rules read at the start of a session, then structure, the brief form whose empty slot shows without anyone remembering to look. Strongest is a deterministic check, a small program firing from outside the planner’s attention at the exact moment a rule applies. The portable kit installs two by default: one watches subagent launches and names any brief slot still unfilled, the other watches the end of a turn that collected results and, finding no outcome named, asks once for it.

Both advise and neither refuses an action, and both fail open, letting the work proceed and logging the failure if something inside them breaks. The outcome nudge costs one extra turn whenever it fires, because the interface it hangs on offers no advisory channel. One environment variable silences both. By this workflow’s own standard the pair is unproven: they behave as specified against their test cases, which makes them fixture-measured, while no live session has yet measured whether they change what the planner does.

What generalizes, and what is specific to your project

You may want to carry this way of working to another project someday, so this section separates the **transferable pattern** from the **local specifics**. The transferable pattern, the schema, is most of what this guide has taught. A coordinator that only decomposes, routes, synthesizes, and decides, while workers do all the domain work, transfers to any supervised pipeline. Code counts alongside them as a third participant of equal rank, and that ranking transfers too.

The same is true of the other core moves: routing deterministic work to code and judgment work to agents; tiering models by difficulty to control cost; choosing an arrangement by how the pieces of work depend on each other; letting the coordinator touch primary sources only to verify a load-bearing claim; refusing to call anything “working” before it is measured; and leaning on an append-only trail and dated backups when there is no version control. These are platform-agnostic ideas. They would serve a data-cleaning cascade or a literature-review pipeline just as well as they serve your project, which develops agents and skills for improved science writing and explanatory writing.

Four parts of the pattern carry honest scope-limits. The rule that *all* inter-agent data flows through files rests entirely on workers being isolated; on a platform where agents share a live blackboard

or a common memory, files remain excellent for provenance but stop being the *only* channel. The rule that a redesign waits for the next launch rests on a subagent not being re-briefable while it runs; where a worker can take a new task mid-run, you can redirect it as it works and there is no need to wait for the collect. And the full, careful brief, persisted to a file and carrying a stuck rule and a report cap, is sized for non-trivial work; a one-line mechanical fan-out does not need the full apparatus. Routing a whole chunk to its own planner rests on a platform where a worker can itself delegate, which this one was tested to do, one level deep and no further. The pattern holds within these limits.

The local specifics are few. The absolute workspace path and the scope-isolation rule come from this sandbox and its history. The model values are this toolkit's own policy: the exact top model, the ban on the next generation, and the ban on unpinned names. The named specialists, the folder layout, the detector script, and the source PDFs belong to this project alone.

How you steer it

Steering this loop comes down to a handful of levers. You set the two things the whole loop is anchored to: the **goal** and the **scope**. Every cycle is re-fitted to the goal at its close, and no subagent may step outside the scope you set without your explicit say-so, so these two inputs shape everything downstream. You read the **durable outputs and the ledgers**, which is where the state of the work lives, in files you can open, not a conversation that scrolls away: the reports each cycle produces, the efficacy ledger that tells you what has been measured versus merely built, and the change log that tells you everything that has happened.

When you step in, where your instruction lands depends on how big it is. A correction to work already under way can reach a running subagent at its next step. A change of direction lands at the next collect, where the planner re-briefs only the pieces of work it touches. The stuck rule runs the other way: instead of waiting for you to intervene, a subagent heading down a wrong path stops and surfaces what it found. Now that you understand how the system works, you are in a better position to predict what it will do and to steer it.